

---

# Question Answering on Wikipedia using Memory Networks

---

Anonymous Author(s)

Affiliation

Address

email

## Abstract

In this project, we investigate the task of building a Question Answering system using deep neural networks augmented with a memory component. Our goal is to implement memory based neural network models, MemN2N and its extensions described in Sukhbaatar et al. (2015)[4], Attentive Reader described by Hermann et al.(2015) [5] and apply it on the wiki-reading QA tasks introduced by Hewlett et al.(2016) [1]. Both the above models have the distinction of being weakly supervised and trained end to end, making it viable for realistic settings. They seem to show promising performance on synthetic question answering datasets and other language task like language modeling. We realize that unlike simulated datasets like bAbI, the memory models are not sufficient to achieve satisfactory performance on real-world QA datasets like Wiki QA. We leverage the works of these authors and explore few extensions to the proposed systems to make it work on these large datasets.

## 1 Introduction

Two interesting challenges in natural language processing research have been to build models capable of multiple computational steps to achieve a task and models that could describe long term dependencies in sequential data. The task of question answering (QA) aptly fits the former. Additionally, QA is extremely broad as many NLP tasks can be reformulated in the QA setup. This implies that devising better models for improving QA can be quite useful in different directions.

Fundamentally, QA systems need to perform two tasks: retrieval and inference. The QA system need to store the knowledge available in some convenient internal representation and subsequently search through the representation to find relevant bits to answer the question. Analyzing the knowledge to derive the required answer requires some form of inference.

In this work we investigate a particular class of learning models called memory networks. Memory networks consists of an inference component combined with a long term memory component. The memory component can be read, written to with the goal of predicting the best possible answer. We use large scale QA task as our benchmark to evaluate the effectiveness of these models.

## 2 Background

### 2.1 Memory Networks

A memory network consists of a memory  $m$  (an array of objects indexed by  $m_i$ ) and four (potentially learned) components I, G, O and R as follows:

I: (input feature map) – converts the incoming input to the internal feature representation

33 G: (generalization) – updates old memories given the new input. We call this generalization as there  
 34 is an opportunity for the network to compress and generalize its memories at this stage for some  
 35 intended future use.

36 O: (output feature map) – produces a new output (in the feature representation space), given the new  
 37 input and the current memory state.

38 R: (response) – converts the output into the response format desired. For example, a textual response  
 39 or an action.

40 Given an input  $x$  (e.g., an input character, word or sentence depending on the granularity chosen, an  
 41 image or an audio signal) the flow of the model is as follows:

- 42 1. Convert  $x$  to an internal feature representation  $I(x)$ .
- 43 2. Update memories  $m_i$  given the new input:  $m_i = G(m_i, I(x), m), \forall i$ .
- 44 3. Compute output features  $o$  given the new input and the memory:  $o = O(I(x), m)$
- 45 4. Finally, decode output features  $o$  to give the final response:  $r = R(o)$

46 I component: Component I can make use of standard pre-processing, e.g., parsing, coreference and  
 47 entity resolution for text inputs. It could also encode the input into an internal feature representation,  
 48 e.g., convert from text to a sparse or dense feature vector.

49 G component: The simplest form of G is to store  $I(x)$  in a “slot” in the memory:

$$m_{H(x)} = I(x) \quad (1)$$

50 where  $H(\cdot)$  is a function selecting the slot. That is, G updates the index  $H(x)$  of  $m$ , but all other  
 51 parts of the memory remain untouched. More sophisticated variants of G could go back and update  
 52 earlier stored memories (potentially, all memories) based on the new evidence from the current input  
 53  $x$ . If the input is at the character or word level one could group inputs (i.e., by segmenting them into  
 54 chunks) and store each chunk in a memory slot.

55 If the memory becomes full, a procedure for “forgetting” could also be implemented by  $H$  as it  
 56 chooses which memory is replaced, e.g.,  $H$  could score the utility of each memory, and overwrite the  
 57 least useful. We have not explored this experimentally yet.

58 O and R components: The O component is typically responsible for reading from memory and  
 59 performing inference, e.g., calculating what are the relevant memories to perform a good response.  
 60 The R component then produces the final response given O. For example in a question answering  
 61 setup O finds relevant memories, and then R produces the actual wording of the answer, e.g., R could  
 62 be an RNN that is conditioned on the output of O.

63 We investigate particular instantiation of memory networks where the components are neural networks.  
 64 We refer to these as memory neural networks (MemNNs) and Attentive Reader neural networks  
 65 (ARNNs).

## 66 2.2 Memory Neural Networks(MemNNs)

67 Our model takes a discrete set of inputs  $x_1, \dots, x_n$  that are to be stored in the memory, a query  $q$ ,  
 68 and outputs an answer  $a$ . Each of the  $x_i, q$ , and  $a$  contains symbols coming from a dictionary with  $V$   
 69 words. The model writes all  $x$  to the memory up to a fixed buffer size, and then finds a continuous  
 70 representation for the  $x$  and  $q$ . The continuous representation is then processed via multiple hops to  
 71 output  $a$ . This allows backpropagation of the error signal through multiple memory accesses back to  
 72 the input during training.

73 Input memory representation: Suppose we are given an input set  $x_1, \dots, x_i$  to be stored in memory.  
 74 The entire set of  $x_i$  are converted into memory vectors  $m_i$  of dimension  $d$  computed by embedding  
 75 each  $x_i$  in a continuous space, in the simplest case, using an embedding matrix  $A(d, V)$ . The query  $q$  is  
 76 also embedded (again, in the simplest case via another embedding matrix  $B$  with the same dimensions  
 77 as  $A$ ) to obtain an internal state  $u$ . In the embedding space, we compute the match between  $u$  and  
 78 each memory  $m_i$  by taking the inner product followed by a softmax:

$$p_i = \text{Softmax}(u^T m_i) \quad (2)$$

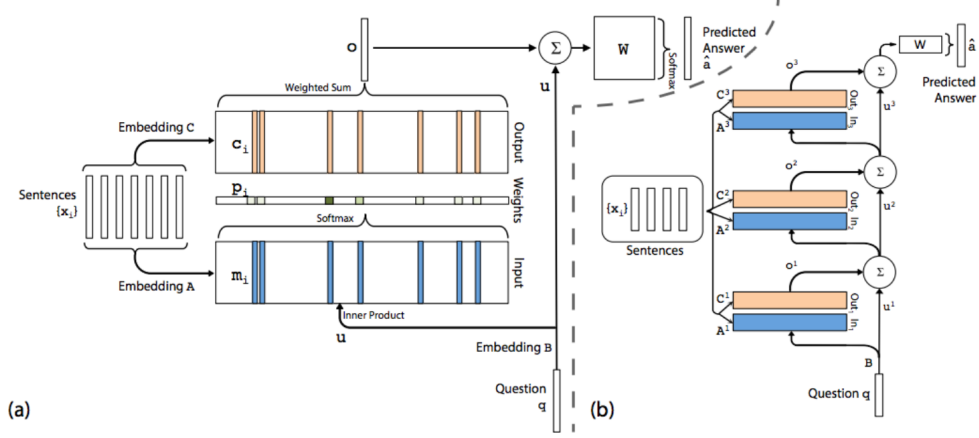


Figure 1: (a): Single layer version. (b): Three layer version

79 where  $Softmax(z_i) = e^{z_i} / \sum_j e^{z_j}$ . Defined in this way  $p$  is a probability vector over the inputs.

80 Output memory representation: Each  $x_i$  has a corresponding output vector  $c_i$  (given in the simplest  
81 case by another embedding matrix  $C$ ). The response vector from the memory  $o$  is then a sum over the  
82 transformed inputs  $c_i$ , weighted by the probability vector from the input:

$$o = \sum_i p_i * c_i \quad (3)$$

83 Because the function from input to output is smooth, we can easily compute gradients and backpropa-  
84 gate through it.

85 Generating the final prediction: In the single layer case, the sum of the output vector  $o$  and the input  
86 embedding  $u$  is then passed through a final weight matrix  $W(V,d)$  and a softmax to produce the  
87 predicted label:

$$\hat{a} = Softmax(W(o + u)) \quad (4)$$

88 During training, all three embedding matrices  $A, B$  and  $C$ , as well as  $W$  are jointly learned by  
89 minimizing a standard cross-entropy loss between  $\hat{a}$  and the true label  $a$ . Training is performed using  
90 stochastic gradient descent.

### 91 2.3 Attentive Reader(ARNNs)

92 The naive LSTM Reader must propagate dependencies over long distances in order to connect queries  
93 to their answers. The fixed width hidden vector forms a bottleneck for this information flow that the  
94 attentive reader circumvents using an attention mechanism. This attention model first encodes the  
95 document and the query using separate bidirectional single layer LSTMs.

96 The first step of the model is to apply an encoding to the context  $c = c_1, c_2, \dots, c_n$  through a  
97 bidirectional LSTM network. The resulting forward and output vectors  $\vec{H}_c = (\vec{h}_c)_i$  and  $\overleftarrow{H}_c = (\overleftarrow{h}_c)_i$   
98 are concatenated into a final paragraph representation  $(H_c)_i = ((\vec{H}_c)_i, (\overleftarrow{H}_c)_i)$ . Each output vector is  
99 chosen to be in  $R^h$ , where  $h = 100$ . The concatenation is thus in  $R^{2h}$ . Another bidirectional LSTM  
100 network is also applied to map the question sequence  $q = q_1, q_2, \dots, q_m$  into  $h_q = ((\vec{h}_q)_m, (\overleftarrow{h}_q)_m)$ .  
101 All of the output states are used in the paragraph representation but only the final states are used in the  
102 question representation. Because the questions are typically shorter in length and can be adequately  
103 represented with less information.

104 The next step is to compare the question and context embeddings so that we may select the most  
105 relevant information. This is done by computing an output vector  $a$ :

$$a_i = Softmax(q^T W_1 (H_c)_i), i = 1, \dots, n \quad (5)$$

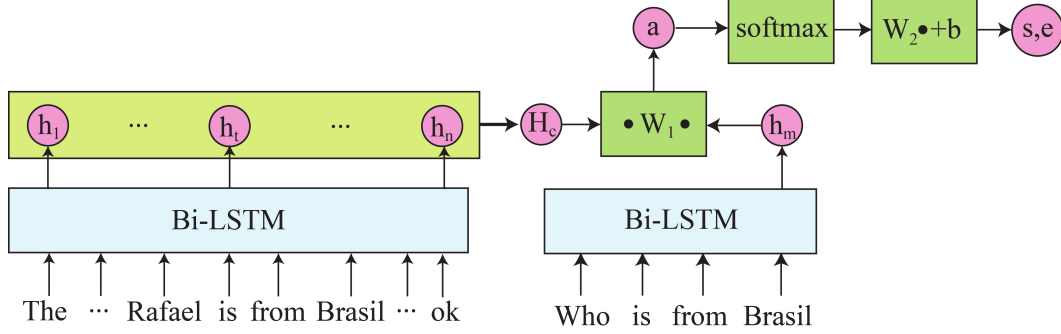


Figure 2: Attentive reader neural network

106 Here we use  $W_1 = R^{4h}$  for a bilinear term which provides more flexibility in computing the similarity  
 107 between  $c$  and  $q$  than a direct product. The softmax is applied to provide nonlinear normalization.

108 This can be viewed as a generalization of the application of memory networks to question answering.  
 109 This model employs an attention mechanism at the sentence level where each sentence is represented  
 110 by a bag of embeddings. The attentive reader employs a fine grained token level attention mechanism  
 111 where the tokens are embedded given their entire future and past context with in the story.

### 112 3 DataSet

113 The WIKI READING task requires to predict textual values from the open knowledge base Wikidata  
 114 given text from the corresponding articles on Wikipedia. Example instances are shown in Table 1,  
 115 illustrating the variety of subject mat- ter and sub-tasks. The dataset contains 18.58M in- stances  
 116 across 884 sub-tasks, split roughly evenly between classification and extraction

	Categorization		Extraction	
Document	Folkart Towers are twin skyscrapers in the Bayrakli district of the Turkish city of Izmir. Reaching a structural height of 200 m (656 ft) above ground level, they are the tallest...	Angeles blancos is a Mexican telenovela produced by Carlos Sotomayor for Televisa in 1990. Jacqueline Andere, Rogelio Guerra and Alfonso Iturralde star as the main...	Canada is a country in the northern part of North America. Its ten provinces and three territories extend from the Atlantic to the Pacific and northward into the Arctic Ocean, . . .	Breaking Bad is an American crime drama television series created and produced by Vince Gilligan. The show originally aired on the AMC network for five seasons, from January 20, 2008, to . . .
Property	country	original language of work	located next to body of water	start time
Answer	Turkey	Spanish	Atlantic Ocean, Arctic Ocean, Pacific Ocean	20 January 2008

Table 1: Examples instances from WIKI READING. The task is to predict the answer given the document and property. Answer tokens that can be extracted are shown in bold, the remaining instances require classification or another form of inference.

117 The dataset contains 4.7M unique Wikipedia ar- ticles, meaning that roughly 80 percent of the English  
 118 language Wikipedia is represented. Multiple in- stances can share the same document, with a mean  
 119 of 5.31 instances per article (median: 4,max:879).The most common categories of documents are  
 120 human,taxon,film,album, and human settlement, making up 48.8 percent of the documents and 9.1  
 121 percent of the instances. The mean and median document lengths are 489.2 and 203 words.

## 122 4 Approach

123 Recently, neural network architectures for QA have been shown to meet or exceed the performance  
 124 of traditional methods. The move to deep neural networks also allows for new ways of combin- ing  
 125 the property and document, inspired by recent research in the field of question answering (with the  
 126 property serving as a question). In memory network models, the question could be used to compute  
 127 a form of attention over the document, to effectively focus the model on the most predictive words  
 128 or phrases (Sukhbaatar et al., 2015; Hermann et al., 2015). As this is currently an ongoing field of  
 129 research,we have implemented them both to compare them on common grounds. We now describe  
 130 these methods and introduce some adjustments to tackle the data set size.

131 **Baseline LSTM** This model is a simplified version of the Deep LSTM Reader proposed by Her-  
 132 mann et al. (2015). In this model, an LSTM reads the property and document sequences word-by-word  
 133 and the final state is used as the joint representation. This is the simplest model that respects the order  
 134 of the words in the document. In our implementation we use a single layer instead of two and a larger  
 135 hidden size.

136 **Memory Networks** Our implementation closely follows the End-to-End Memory Network pro-  
 137 posed in Sukhbaatar et al. (2015). This model maps a property  $p$  and a list of sentences  $x_1, \dots, x_n$  to  
 138 a joint representation  $y$  by attending over sentences in the document as follows: The input encoder  $I$   
 139 converts a sequence of words  $x_i = (x_{i1}, \dots, x_{iL_i})$  into a vector using an embedding matrix (equation  
 140 6), where  $L_i$  is the length of sentence  $i$ .The property is encoded with the embedding matrix  $U$  (eqn.  
 141 7). Each sentence is encoded into two vectors, a memory vector (eqn. 7) and an output vector (eqn.  
 142 9), with embedding matrices  $M$  and  $C$ , respectively. The property encoding is used to compute a  
 143 normalized attention vector over the memories (eqn. 10). The joint representation is the sum of the  
 144 output vectors weighted by this attention

$$I(x_i, W) = \sum_j W x_{ij} \quad (6)$$

$$m_i = I(x_i, M) \quad (7)$$

$$c_i = I(x_i, C) \quad (8)$$

$$p_i = Softmax(q^T m_i) \quad (9)$$

$$y = u + \sum_i p_i c_i \quad (10)$$

149 **Attentive Reader** This model, also presented in Hermann et al. (2015), uses an attention mecha-  
 150 nism to better focus on the relevant part of the document for a given property. Specifically, Attentive  
 151 Reader first generates a representation  $u$  of the property using the final state of an LSTM while  
 152 a second LSTM is used to read the document and generate a representation  $z_t$  for each word.  
 153 Then,conditioned on the property encoding  $u$ , a normalized attention is computed over the document  
 154 to produce a weighted average of the word representations  $z_t$ , which is then used to generate the joint  
 155 representation  $y$ . More precisely:

$$m_t = \tanh(W_1 \text{concat}(z_t, u)) \quad (11)$$

$$\alpha_t = \exp(v^T m_t) \quad (12)$$

$$r = \sum_t \frac{\alpha_t}{\sum_i \alpha_i} z_t \quad (13)$$

$$y = \tanh(W_2 \text{concat}(r, u)) \quad (14)$$

## 5 Training

We used keras as the deep learning framework on top of tensorflow. The hardware we used is the NVIDIA TITAN x GPU that has 12 GB of memory.

Method	Embedding Dimension	Vocab Size	Batch Size	Learnable Parameters
Baseline LSTM	64	200K	100	38M
MNN	300	200k	1000	91M
ARNN	128	200K	500	46M

Table 2: Training Details

Pruning: We weren't able to successfully train the model against some story sizes, being the entire wikipedia pages. So we pruned the document into manageable chunks by computing the similarity between the sentences and question and picking the top k. Clearly this isn't an ideal way, but we weren't able to investigate the model with some certainty with vocab size greater than this.

Evaluation: We evaluated all three methods on randomly sampled batches from separate test set using a single scoring framework. An answer is correct when there is an exact string match between the predicted answer and the gold answer. However, some answers are composed from a set of values. To handle this, we define the Mean F1 score as follows: For each instance, we compute the F1-score (harmonic mean of precision and recall) as a measure of the degree of overlap between the predicted answer set and the gold set for a given instance. The resulting per- instance F1 scores are then averaged to produce a single dataset-level score. This allows a method to obtain partial credit for answers.

Some of the date related question types are clearly hard no matter which model we use. It seems the number of supporting facts involved, the complexity of the relation itself, and the length of the context makes it hard for the models to reason about.

## GitHub

Project Repository

## Results

Question Type	Baseline	MNN	ARNN
"country"	0.14	0.46	0.38
"occupation"	0.19	0.37	0.35
"genre"	0.17	0.32	0.31
"sport"	0.11	0.41	0.43
"date"	0.08	0.11	0.07

Table 3: Accuracy on hold out data

## Conclusion and Future Work

In this project, we tried to evaluate the memory based neural models for the task of question answering from raw wikipedia pages. We started with a naive LSTM, based on the similarity of the question and input facts. The accuracy of the baseline is below .2 for all categories of the question. Then our model of choices MNN's and ARNN's did much better than the naive LSTM model in relatively better training time. We haven't tried multi layers in both the models, particularly MNN's.

We have pruned sentences to contain data set with in the available memory footprint, which turned out to be the bottleneck, rather training a component separately to spot the relevant part of the story based on supervision seems promising lead. Still the training of the model on complete data set is possible if we could parallelize it. We haven't investigated much on that. Our best model was able to achieve 0.43 accuracy on the hold out set. At the moment, the models are trained separately based on question type, generalising it would be another challenge as well.




Figure 3: Loss (MNN)

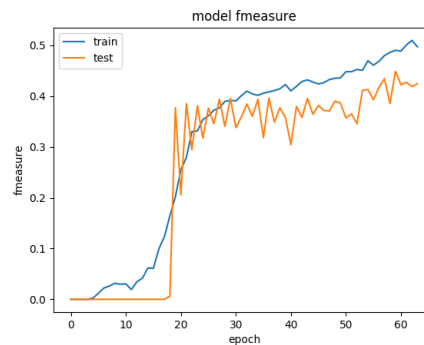


Figure 4: F1 (MNN)

## References

- [1] Daniel Hewlett., Alexandre Lacoste., Llion Jones., Illia Polosukhin., & Andrew Fandrianto (2016) WikiReading: A Novel Large-scale Language Understanding Task over Wikipedia, *Proceedings of the 54th Annual Meeting of the Association for Computational Linguistics* 7, pp. 1535-1545.
- [2] Jason Weston., Sumit Chopra. & Antoine Bordes (2014) Memory Networks.
- [3] Antoine Bordes, Nicolas Usunier, Sumit Chopra, & Jason Weston. (2015) Large-scale Simple Question Answering with Memory Networks.
- [4] Sainbayar Sukhbaatar, Arthur Szlam, Jason Weston, & Et Al. (2015) End-To-End Memory Networks.
- [5] Karl Moritz Hermann, Tomas Kocisky, Edward Grefenstette, & Et Al. (2015) Teaching Machines to Read and Comprehend.
- [6] Ankit Kumar, Ozan Irsoy, Peter Ondruska, & Et Al. (2016) Ask Me Anything: Dynamic Memory Networks for Natural Language Processing, *Proceedings of The 33rd International Conference on Machine Learning* , PMLR 48:1378-1387.
- [7] Matt J. Kusner, Yu Sun, Nicholas I. Kolkin, & Kilian Q. Weinberger (2015) From word embeddings to document distances, *Proceeding ICML'15 Proceedings of the 32nd International Conference on International Conference on Machine Learning* , Volume 37 Pages 957-966.
- [8] Matthew Richardson, Christopher J.C. Burges, & Erin Renshaw (2013) MCTest: A Challenge Dataset for the Open-Domain Machine Comprehension of Text, *Proceedings of the 2013 Conference on Empirical Methods in Natural Language Processing (EMNLP 2013)*
- [9] Jason Weston, Antoine Bordes, & Et Al (2015) Towards AI-Complete Question Answering: A Set of Prerequisite Toy Tasks, *Proceedings of the 2016 International Conference on Learning Representations (ICLR 2016)*